


How to Shadow Every Byte of Memory Used by a Program



Nicholas Nethercote – National ICT Australia


Julian Seward – OpenWorks LLP



Shadow memory tools



- Shadow every byte of memory with another value that describes it



- This talk:
 - Why shadow memory is useful
 - How to implement it well



Examples




	Tool(s)	Shadow memory helps find...
bugs	Memcheck , Purify	Memory errors
	Eraser, DRD, Helgrind, etc.	Data races
	Hobbes	Run-time type errors
	Annelid	Array bounds violations
security	TaintCheck, LIFT, TaintTrace “Secret tracker”	Uses of untrusted values Leaked secrets
properties	Redux DynCompB pinSEL	Dynamic dataflow graphs Invariants System call side-effects



Shadow memory is difficult



- Performance
 - Lots of extra state, many operations instrumented
- Robustness



- Trade-offs must be made

An example tool: Memcheck






Memcheck



- Three kinds of information:
 - A (“addressability”) bits: 1 bit / memory byte
 - V (“validity”) bits: 1 bit / register bit, 1 bit / memory bit
 - Heap blocks: location, size, allocation function
- Memory information:



V bits only used if A bit is “addressable”

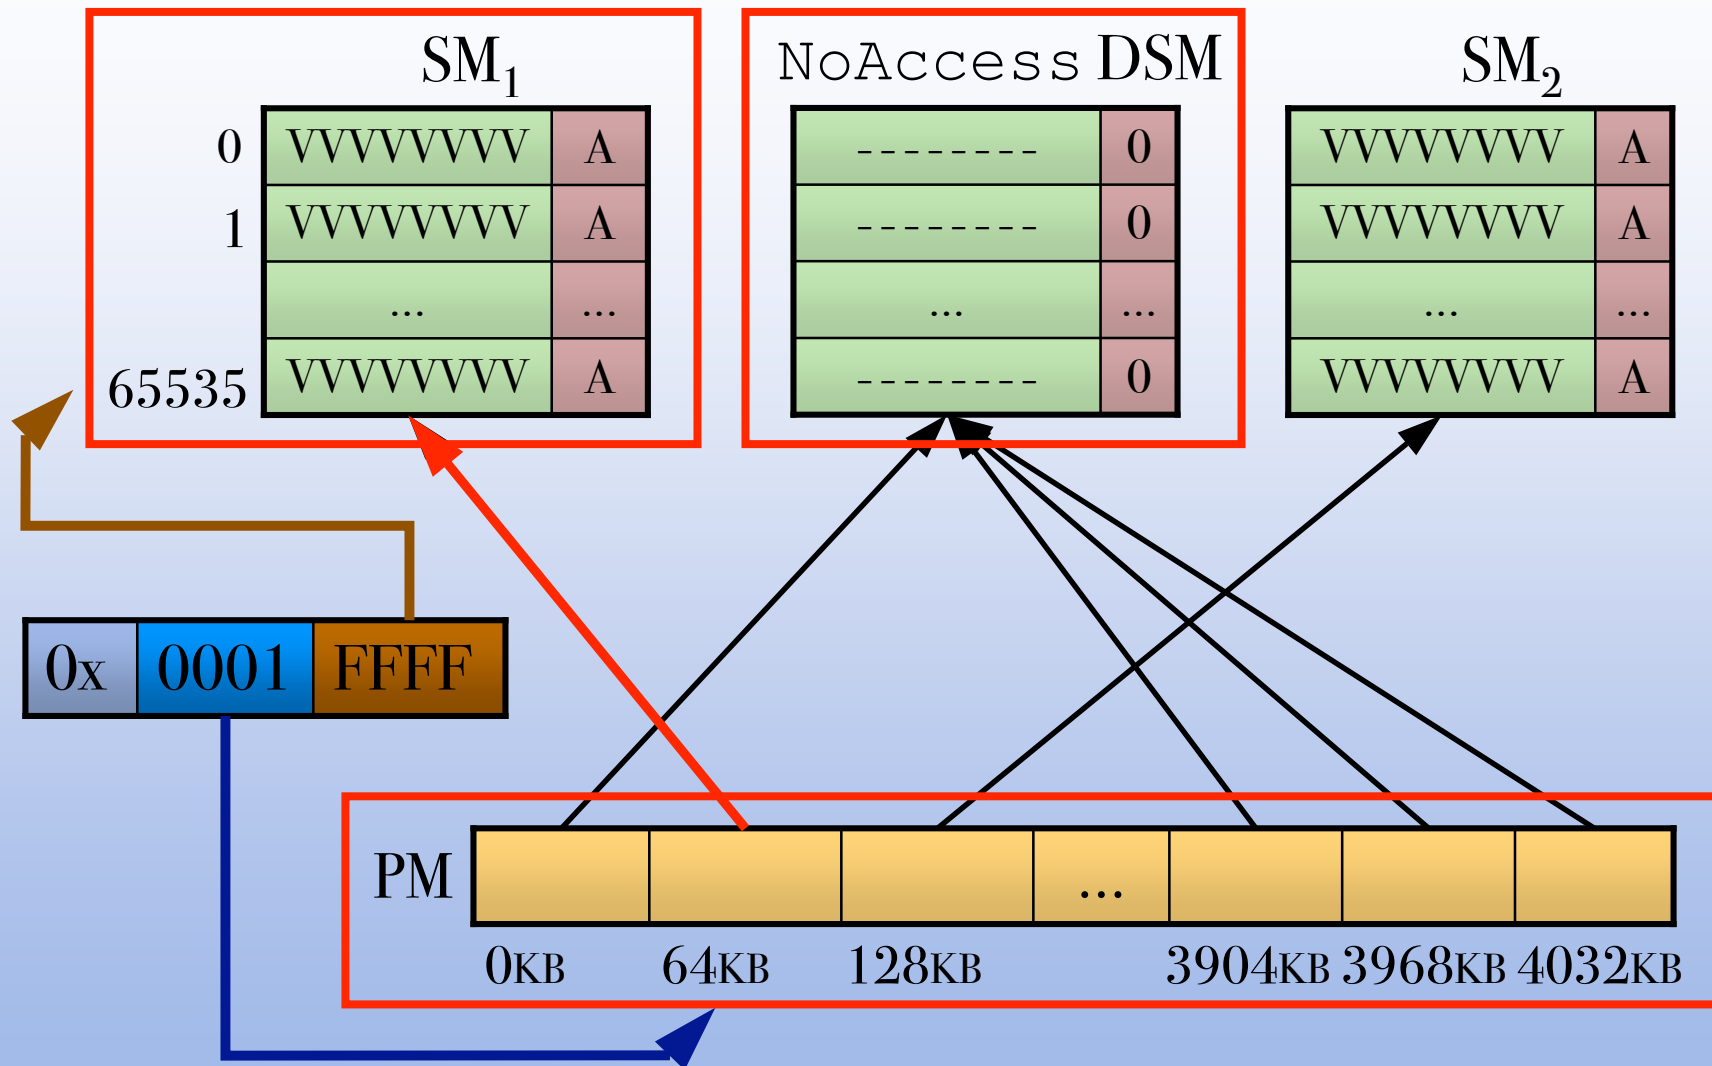


A simple implementation





Basics (I)





Basics (II)



- Multi-byte shadow accesses:
 - Combine multiple single-byte accesses
 - Complain if any unaddressable bytes accessed
 - Values loaded from unaddressable bytes marked as defined
- Range-setting (`set_range`)
 - Loop over many bytes, one at a time
- Range-checking
 - E.g.: `write(fd, buf, n)` -- check `n` bytes in `buf`
- Slow-down: 209.6x



Complications



- Corruption of shadow memory
 - Possible with a buggy program
 - Originally used x86 segmentation, but not portable
 - Keep original and shadow memory far apart, and pray
- 64-bit machines
 - Three- or four-level structure would be slow
 - Two level structure extended to handle 32GB
 - Slow auxiliary table for memory beyond 32GB
 - Better solution is an open research question

Four optimisations





#1: Faster loads and stores



- Multi-byte loads/stores are very common
 - N separate lookups accesses is silly (where N = 2, 4, or 8)
- If access is aligned, fully addressable
 - Extract/write V bits for N shadow bytes at once
 - Else fall back to slow case: 1 in a 1000 or less
- Slow-down: 56.2x
 - 3.73x faster



#2: Faster range-setting



- Range-setting large areas is common
 - Vectorise `set_range`
 - 8-byte stride works well
- Replacing whole SMs
 - If marking a 64KB chunk as `NoAccess`, replace the SM with the `NoAccess DSM`
 - Add `Defined` and `Undefined DSMs`
 - Large read-only code sections covered by `Defined DSM`
- Slow-down: 34.7x
 - 1.62x faster, 1.97x smaller



#3: Faster SP updates



- Stack pointer (SP) updates are very common
- Inc/dec size often small, statically known
 - E.g. 4, 8, 12, 16, 32 bytes
- More specialised range-setting functions
 - Unrolled versions of `set_range ()`
- Slow-down: 27.2x
 - 1.28x faster



#4: Compressed V bits



- Partially-defined bytes (PDBs) are rare
 - Memory: 1 A bit + 8 V bits \rightarrow 2 VA bits
 - Four states: NoAccess, Undefined, Defined, PartDefined
 - Full V bits for PDBs in secondary V bits table
 - Registers unchanged – still 8 V bits per byte
- Slow-down: 23.4x
 - 4.29x smaller, 1.16x faster
- Obvious in hindsight, but took 3 years to identify



Discussion



- Optimising principles:
 - Start with a simple implementation
 - Make the common cases fast
 - Exploit redundancy to reduce data sizes
- Novelty?
 - First detailed description of Memcheck's shadow memory
 - First detailed description of a two-level table version
 - First detailed evaluation of shadow memory
 - Compressed V bits

Evaluation





Robustness



- Two-level table is very flexible
 - Small shadow memory chunks, each can go anywhere
- Earlier versions required large contiguous regions
 - Some programs require access to upper address space
 - Some Linux kernels have trouble mmap'ing large regions
 - Big problems with Mac OS X, AIX, other OSes
- Memcheck is robust
 - Standard Linux C and C++ development tool
 - Official: Linux, AIX; experimental: Mac OS X, FreeBSD



SPEC 2000 Performance



Tool	Slow-down	Relative improvement
No instrumentation	4.3x	
Simple Memcheck	209.6x	
+ faster loads/stores	56.2x	3.73x faster
+ faster range-setting	34.7x	1.62x faster, 1.97x smaller
+ faster SP updates	27.2x	1.28x faster
+ compressed V bits	23.4x	1.16x faster, 4.29x smaller
Overall improvement		8.9x faster, 8.5x smaller

- Shadow memory causes about half of Memcheck's overhead



Performance observations




- Performance is a traditional research obsession
 - “The *subjective* issues are important – ease of use and robustness, but performance is the item which would be most interesting for the audience.” (my emphasis)
- Users: slowness is #1 survey complaint
 - But most user emails are about bugs or interpreting results
 - Zero preparation is a big win
- Cost/benefit
 - People will use slow tools if they are sufficiently useful



Alternative implementation



- “Half-and-half”
 - Used by Hobbes, TaintTrace, (with variation) LIFT



- Compared to two-level table
 - Faster
 - Not robust enough for our purposes

If you remember nothing else...





Take-home messages



- Shadow memory is powerful
- Shadow memory can be implemented well
- Implementations require trade-offs

